

# Join Statements

## Introduction

In this section, you will learn about several types of **JOIN** statements. Joins are the primary mechanism for combining data from multiple tables. In order to do this, you define the common attribute(s) between tables in order for them to be combined.

## Objectives

You will be able to:

- Write SQL queries that make use of various types of joins
- Compare and contrast the various types of joins
- Discuss how primary and foreign keys are used in SQL
- Decide and perform whichever type of join is best for retrieving desired data

## CRM ERD

In almost all industry cases, rather than just working with a single table you will generally need data from multiple tables. Doing this requires the use of **joins** using shared columns from the two tables. For example, here's a diagram of a mock customer relationship management (CRM) database.

## Connecting to the Database

As usual, you'll start by connecting to the database.

In the cell below, type the code to import `sqlite` and `pandas` with the standard alias. Then in the next cell create a connection to the database `data.sqlite` and assign it to a variable:

```
# replace this comment with the code to import the libraries
import sqlite3
import pandas as pd

# replace this comment with the code to create a connection to the
database data.database
conn = sqlite3.connect('data.sqlite')

import sqlite3
import pandas as pd
conn = sqlite3.connect('data.sqlite')
```

## Displaying Product Details Along with Order Details

Let's say you need to generate a report that includes details about products from orders. To do that, we would need to take data from multiple tables in a single statement. To do this we will use **JOIN**.

In the cell below, type the query to select all records from **orderdetails** and **products** and join them using their common key **productCode** and display the first 10.

```
# replace None with the query to join orderdetails and products on productCode
```

```
query = '''
```

```
    SELECT * FROM orderdetails
```

```
    JOIN products
```

```
    ON orderdetails.productCode = products.productCode
```

```
    LIMIT 10;
```

```
'''
```

```
pd.read_sql(query, conn)
```

	orderNumber	productCode	quantityOrdered	priceEach
orderLineNumber \				
0	10100	S18_1749	30	136.00
3				
1	10100	S18_2248	50	55.09
2				
2	10100	S18_4409	22	75.46
4				
3	10100	S24_3969	49	35.29
1				
4	10101	S18_2325	25	108.06
4				
5	10101	S18_2795	26	167.06
1				
6	10101	S24_1937	45	32.53
3				
7	10101	S24_2022	46	44.35
2				
8	10102	S18_1342	39	95.55
2				
9	10102	S18_1367	41	43.13
1				

	productCode	productName	productLine
\			
0	S18_1749	1917 Grand Touring Sedan	Vintage Cars
1	S18_2248	1911 Ford Town Car	Vintage Cars
2	S18_4409	1932 Alfa Romeo 8C2300 Spider Sport	Vintage Cars

3	S24_3969	1936 Mercedes Benz 500k Roadster	Vintage Cars
4	S18_2325	1932 Model A Ford J-Coupe	Vintage Cars
5	S18_2795	1928 Mercedes-Benz SSK	Vintage Cars
6	S24_1937	1939 Chevrolet Deluxe Coupe	Vintage Cars
7	S24_2022	1938 Cadillac V-16 Presidential Limousine	Vintage Cars
8	S18_1342	1937 Lincoln Berline	Vintage Cars
9	S18_1367	1936 Mercedes-Benz 500K Special Roadster	Vintage Cars

	productScale	productVendor	\
0	1:18	Welly Diecast Productions	
1	1:18	Motor City Art Classics	
2	1:18	Exoto Designs	
3	1:24	Red Start Diecast	
4	1:18	Autoart Studio Design	
5	1:18	Gearbox Collectibles	
6	1:24	Motor City Art Classics	
7	1:24	Classic Metal Creations	
8	1:18	Motor City Art Classics	
9	1:18	Studio M Art Models	

	productDescription	quantityInStock
\		
0	This 1:18 scale replica of the 1917 Grand Tour...	2724
1	Features opening hood, opening doors, opening ...	540
2	This 1:18 scale precision die cast replica fea...	6553
3	This model features grille-mounted chrome horn...	2081
4	This model features grille-mounted chrome horn...	9354
5	This 1:18 replica features grille-mounted chro...	548
6	This 1:24 scale die-cast replica of the 1939 C...	7332
7	This 1:24 scale precision die cast replica of ...	2847
8	Features opening engine cover, doors, trunk, a...	8693
9	This 1:18 scale replica is constructed of heav...	8635

	buyPrice	MSRP
0	86.70	170.00
1	33.30	60.54
2	43.26	92.03
3	21.75	41.03
4	58.48	127.13
5	72.56	168.75
6	22.57	33.19
7	20.61	44.80
8	60.62	102.74
9	24.26	53.91

---

Expected Output

---

## Compared to the Individual Tables:

### orderdetails Table:

In the cell below, type the code to select all records from `orderdetails` and display the first 10

```
# replace None with the query to display the first 10 records in orderdetails
query = """SELECT * FROM orderdetails LIMIT 10"""
pd.read_sql(query, conn)
```

	orderNumber	productCode	quantityOrdered	priceEach
orderLineNumber				
0	10100	S18_1749	30	136.00
3				
1	10100	S18_2248	50	55.09
2				
2	10100	S18_4409	22	75.46
4				
3	10100	S24_3969	49	35.29
1				
4	10101	S18_2325	25	108.06
4				
5	10101	S18_2795	26	167.06
1				
6	10101	S24_1937	45	32.53
3				
7	10101	S24_2022	46	44.35
2				
8	10102	S18_1342	39	95.55
2				

9	10102	S18_1367	41	43.13
1				

---

## Expected Output

---

### products Table:

In the cell below, type the code to select all records from `products` and display the first 10

```
# replace None with the query to display the first 10 records in products
```

```
query = None
pd.read_sql(query, conn)
```

```
-----
-----
ValueError                                Traceback (most recent call
last)
c:\Users\rurig\anaconda3\envs\learn-env\lib\site-packages\pandas\io\
sql.py in execute(self, *args, **kwargs)
    1680         try:
-> 1681             cur.execute(*args, **kwargs)
    1682         return cur
```

ValueError: operation parameter must be str

The above exception was the direct cause of the following exception:

```
DatabaseError                            Traceback (most recent call
last)
<ipython-input-29-12c934d1a0ed> in <module>
      1 # replace None with the query to display the first 10 records
in products
      2 query = None
----> 3 pd.read_sql(query, conn)

c:\Users\rurig\anaconda3\envs\learn-env\lib\site-packages\pandas\io\
sql.py in read_sql(sql, con, index_col, coerce_float, params,
parse_dates, columns, chunksize)
    481
    482     if isinstance(pandas_sql, SQLiteDatabase):
--> 483         return pandas_sql.read_query(
    484             sql,
    485             index_col=index_col,
```

```
c:\Users\rurig\anaconda3\envs\learn-env\lib\site-packages\pandas\io\
```

```

sql.py in read_query(self, sql, index_col, coerce_float, params,
parse_dates, chunksize)
1725
1726         args = _convert_params(sql, params)
-> 1727         cursor = self.execute(*args)
1728         columns = [col_desc[0] for col_desc in
cursor.description]
1729

c:\Users\rurig\anaconda3\envs\learn-env\lib\site-packages\pandas\io\
sql.py in execute(self, *args, **kwargs)
1691
1692         ex = DatabaseError(f"Execution failed on sql
'{args[0]}': {exc}")
-> 1693         raise ex from exc
1694
1695     @staticmethod

DatabaseError: Execution failed on sql 'None': operation parameter
must be str

```

---

## Expected Output

---

## The USING clause

A more concise way to join the tables, if the column name is identical, is the `USING` clause. Rather than saying `ON tableA.column = tableB.column` we can simply say `USING(column)`. Again, this only works if the column is **identically named** for both tables.

In the cell below, type the query to select all records in `orderdetails` and `products` and join them on `productCode` with the `USING()` clause, and return the first 10 records:

```

# replace None with the query to join orderdetails and proucts on
productCode with the using() clause
query = """
    SELECT * FROM orderdetails
    JOIN products USING(productCode)
    LIMIT 10;
"""
pd.read_sql(query, conn)

```

	orderNumber	productCode	quantityOrdered	priceEach
orderLineNumber \				
0	10100	S18_1749	30	136.00
3				

1	10100	S18_2248	50	55.09
2				
2	10100	S18_4409	22	75.46
4				
3	10100	S24_3969	49	35.29
1				
4	10101	S18_2325	25	108.06
4				
5	10101	S18_2795	26	167.06
1				
6	10101	S24_1937	45	32.53
3				
7	10101	S24_2022	46	44.35
2				
8	10102	S18_1342	39	95.55
2				
9	10102	S18_1367	41	43.13
1				

	productName	productLine
productScale \		
0	1917 Grand Touring Sedan	Vintage Cars
1:18		
1	1911 Ford Town Car	Vintage Cars
1:18		
2	1932 Alfa Romeo 8C2300 Spider Sport	Vintage Cars
1:18		
3	1936 Mercedes Benz 500k Roadster	Vintage Cars
1:24		
4	1932 Model A Ford J-Coupe	Vintage Cars
1:18		
5	1928 Mercedes-Benz SSK	Vintage Cars
1:18		
6	1939 Chevrolet Deluxe Coupe	Vintage Cars
1:24		
7	1938 Cadillac V-16 Presidential Limousine	Vintage Cars
1:24		
8	1937 Lincoln Berline	Vintage Cars
1:18		
9	1936 Mercedes-Benz 500K Special Roadster	Vintage Cars
1:18		

	productVendor \
0	Welly Diecast Productions
1	Motor City Art Classics
2	Exoto Designs
3	Red Start Diecast
4	Autoart Studio Design
5	Gearbox Collectibles

```

6 Motor City Art Classics
7 Classic Metal Creations
8 Motor City Art Classics
9 Studio M Art Models

```

```

                                productDescription  quantityInStock
\
0 This 1:18 scale replica of the 1917 Grand Tour...          2724
1 Features opening hood, opening doors, opening ...          540
2 This 1:18 scale precision die cast replica fea...          6553
3 This model features grille-mounted chrome horn...          2081
4 This model features grille-mounted chrome horn...          9354
5 This 1:18 replica features grille-mounted chro...          548
6 This 1:24 scale die-cast replica of the 1939 C...          7332
7 This 1:24 scale precision die cast replica of ...          2847
8 Features opening engine cover, doors, trunk, a...          8693
9 This 1:18 scale replica is constructed of heav...          8635

    buyPrice    MSRP
0      86.70    170.00
1      33.30     60.54
2      43.26     92.03
3      21.75     41.03
4      58.48    127.13
5      72.56    168.75
6      22.57     33.19
7      20.61     44.80
8      60.62    102.74
9      24.26     53.91

```

---

Expected Output

---

## More Aliasing

You can also assign tables an **alias** by entering an alternative shorthand name. This is slightly different than the previous lesson where we introduced aliases for column names, since now we are aliasing *tables*.



When aliasing columns the goal is usually to improve readability by giving something a more specific or easier-to-read name. For example, `name AS employee_name`, `AVG(AVG) AS average_batting_average`, or `COUNT(*) AS num_products`.

When aliasing tables the goal is usually to shorten the name, in order to shorten the overall query. So typically you'll see examples that alias a longer table name to a one-character or two-character shorthand. For example, `orderdetails AS od` or `products AS p`. (It is also possible to use aliases to clarify what exactly is in a table, like how aliases are used for columns, just less common.)

The following query produces the same result as the previous ones, using aliases `od` and `p` for `orderdetails` and `products`, respectively:

In the following cell, type the following code to demonstrate the use of aliasing:

```
# replace None with the query to demonstrate aliasing
query = """
    SELECT * FROM orderdetails AS od
    JOIN products AS p
    ON od.productCode = p.productCode
    LIMIT 10;
"""
pd.read_sql(query, conn)
```

	orderNumber	productCode	quantityOrdered	priceEach
orderLineNumber \				
0	10100	S18_1749	30	136.00
3				
1	10100	S18_2248	50	55.09
2				
2	10100	S18_4409	22	75.46
4				
3	10100	S24_3969	49	35.29
1				
4	10101	S18_2325	25	108.06
4				
5	10101	S18_2795	26	167.06
1				
6	10101	S24_1937	45	32.53
3				
7	10101	S24_2022	46	44.35
2				
8	10102	S18_1342	39	95.55
2				
9	10102	S18_1367	41	43.13
1				

	productCode	productName	productLine
\			
0	S18_1749	1917 Grand Touring Sedan	Vintage Cars

1	S18_2248	1911 Ford Town Car	Vintage Cars
2	S18_4409	1932 Alfa Romeo 8C2300 Spider Sport	Vintage Cars
3	S24_3969	1936 Mercedes Benz 500k Roadster	Vintage Cars
4	S18_2325	1932 Model A Ford J-Coupe	Vintage Cars
5	S18_2795	1928 Mercedes-Benz SSK	Vintage Cars
6	S24_1937	1939 Chevrolet Deluxe Coupe	Vintage Cars
7	S24_2022	1938 Cadillac V-16 Presidential Limousine	Vintage Cars
8	S18_1342	1937 Lincoln Berline	Vintage Cars
9	S18_1367	1936 Mercedes-Benz 500K Special Roadster	Vintage Cars
	productScale	productVendor	\
0	1:18	Welly Diecast Productions	
1	1:18	Motor City Art Classics	
2	1:18	Exoto Designs	
3	1:24	Red Start Diecast	
4	1:18	Autoart Studio Design	
5	1:18	Gearbox Collectibles	
6	1:24	Motor City Art Classics	
7	1:24	Classic Metal Creations	
8	1:18	Motor City Art Classics	
9	1:18	Studio M Art Models	
		productDescription	quantityInStock
\			
0	This 1:18 scale replica of the 1917 Grand Tour...		2724
1	Features opening hood, opening doors, opening ...		540
2	This 1:18 scale precision die cast replica fea...		6553
3	This model features grille-mounted chrome horn...		2081
4	This model features grille-mounted chrome horn...		9354
5	This 1:18 replica features grille-mounted chro...		548
6	This 1:24 scale die-cast replica of the 1939 C...		7332
7	This 1:24 scale precision die cast replica of ...		2847
8	Features opening engine cover, doors, trunk, a...		8693

9 This 1:18 scale replica is constructed of heav...

8635

	buyPrice	MSRP
0	86.70	170.00
1	33.30	60.54
2	43.26	92.03
3	21.75	41.03
4	58.48	127.13
5	72.56	168.75
6	22.57	33.19
7	20.61	44.80
8	60.62	102.74
9	24.26	53.91

---

## Expected Output

---

Note that just like with column aliases, the **AS** keyword is optional in SQLite. So, instead of **FROM orderdetails AS od** you could write **FROM orderdetails od** with the same outcome.

It is somewhat more common to see **AS** used with column aliases and skipped with table aliases, but again, you'll want to check the syntax rules of your particular type of SQL as well as style guidelines from your employer to know which syntax to use in a professional setting.

## LEFT JOINS

By default a **JOIN** is an **INNER JOIN**, or the intersection between two tables. In other words, the **JOIN** between orders and products is only for **productCodes** that are in both the **orderdetails** and **products** tables. If a product had yet to be ordered (and wasn't in the **orderdetails** table) then it would also not be in the result of the **JOIN**.

The **LEFT JOIN** keyword returns all records from the left table (table1), and the matched records from the right table (table2). The result is NULL from the right side if there is no match.

There are many other types of joins, displayed below. Of these, SQLite does not support outer joins, but it is good to be aware of as more powerful versions of SQL such as PostgreSQL support these additional functions.

For example, the statement

```
SELECT * FROM products LEFT JOIN orderdetails
```

would return all products, even those that hadn't been ordered. You can imagine that all products in inventory should have a description in the product table, but perhaps not every product is represented in the orderdetails table.

In the cell below, type the query to select all records from `products` and join them with all records in `orderdetails` on `productcode` using `LEFT JOIN`, then execute the query and store it in a dataframe named `df`:

```
# replace this comment with the code to create the specified query
q = """SELECT * FROM products LEFT JOIN orderdetails
USING(productCode)"""
# replace this comment with the code to execute the query and store it
in a dataframe named df
df = pd.read_sql(q, conn)
print("Number of records returned:", len(df))
print("Number of records where order details are null:",
len(df[df.orderNumber.isnull()]))
```

Number of records returned: 2997  
Number of records where order details are null: 1

---

## Expected Output

---

Let's take a look at the one record that has null values in the order details:

```
# run this cell with no changes to view the one record with null
values
df[df.orderNumber.isnull()]
```

	productCode	productName	productLine	productScale	\
1122	S18_3233	1985 Toyota Supra	Classic Cars	1:18	

  

	productVendor	\
1122	Highway 66 Mini Classics	

  

	productDescription
1122	This model features soft rubber tires, working...

  

	quantityInStock	\
1122	7733	

  

	buyPrice	MSRP	orderNumber	quantityOrdered	priceEach	\
1122	57.01	107.57	NaN	NaN	NaN	

  

	orderLineNumber
1122	NaN

---

## Expected Output

---

As you can see, it's a rare occurrence, but there is one product that has yet to be ordered.

## Primary Versus Foreign Keys

Another important consideration when performing joins is to think more about the key or column you are joining on. As you'll see in upcoming lessons, this can lead to interesting behavior if the join value is not unique in one or both of the tables. In all of the above examples, you joined two tables using the **primary key**. The primary key(s) of a table are those column(s) which uniquely identify a row. You'll also see this designated in our schema diagram with the asterisk (\*).

You can also join tables using **foreign keys** which are not the primary key for that particular table, but rather another table. For example, `employeeNumber` is the primary key for the `employees` table and corresponds to the `salesRepEmployeeNumber` of the `customers` table. In the `customers` table, `salesRepEmployeeNumber` is only a foreign key, and is unlikely to be a unique identifier, as it is likely that an employee serves multiple customers. As such, in the resulting view `employeeNumber` would no longer be a unique field.

In the cell below, type the query to join `customers` using the alias `c` with `employees` using the alias `e` on the foreign keys `salesRepEmployeeNumber` and `employeeNumber` and order the result by `employeeNumber`, then type the code to execute the query:

```
# replace None with the query to select the desired records
q = """
SELECT * FROM customers AS c
JOIN employees AS e
ON c.salesRepEmployeeNumber = e.employeeNumber
ORDER BY employeeNumber;
"""
# replace this comment with the code to execute the query
pd.read_sql(q, conn)
```

	customerNumber	customerName	contactLastName	\
0	124	Mini Gifts Distributors Ltd.	Nelson	
1	129	Mini Wheels Co.	Murphy	
2	161	Technics Stores Inc.	Hashimoto	
3	321	Corporate Gift Ideas Co.	Brown	
4	450	The Sharp Gifts Warehouse	Frick	
..	...		...	
95	298	Vida Sport, Ltd	Holz	
96	344	CAF Imports	Fernandez	
97	376	Precious Collectables	Urs	
98	458	Corrida Auto Replicas, Ltd	Sommer	
99	484	Iberia Gift Imports, Corp.	Roel	

  

	contactFirstName	phone	addressLine1	\
0	Susan	4155551450	5677 Strong St.	
1	Julie	6505555787	5557 North Pendale Street	
2	Juri	6505556809	9408 Furth Circle	

3	Julie	6505551386	7734 Strong St.
4	Sue	4085553659	3086 Ingle Ln.
..	...	...	...
95	Mihael	0897-034555	Grenzacherweg 237
96	Jesus	+34 913 728 555	Merchants House
97	Braun	0452-076555	Hauptstr. 29
98	Martín	(91) 555 22 82	C/ Araquil, 67
99	José Pedro	(95) 555 82 82	C/ Romero, 33

	addressLine2	city	state	postalCode	...	\
0		San Rafael	CA	97562	...	
1		San Francisco	CA	94217	...	
2		Burlingame	CA	94217	...	
3		San Francisco	CA	94217	...	
4		San Jose	CA	94217	...	
..	...	...	...	...	...	
95		Genève		1203	...	
96	27-30 Merchant's Quay	Madrid		28023	...	
97		Bern		3012	...	
98		Madrid		28023	...	
99		Sevilla		41101	...	

	salesRepEmployeeNumber	creditLimit	employeeNumber	lastName
firstName \				
0	1165	210500	1165	Jennings
Leslie				
1	1165	64600	1165	Jennings
Leslie				
2	1165	84600	1165	Jennings
Leslie				
3	1165	105000	1165	Jennings
Leslie				
4	1165	77600	1165	Jennings
Leslie				
..	...	...	...	...
...				
95	1702	141300	1702	Gerard
Martin				
96	1702	59600	1702	Gerard
Martin				
97	1702	0	1702	Gerard
Martin				
98	1702	104600	1702	Gerard
Martin				
99	1702	65700	1702	Gerard
Martin				

	extension	email	officeCode	reportsTo
jobTitle				
0	x3291	ljennings@classicmodelcars.com	1	1143

```

Sales Rep
1      x3291  ljennings@classicmodelcars.com      1      1143
Sales Rep
2      x3291  ljennings@classicmodelcars.com      1      1143
Sales Rep
3      x3291  ljennings@classicmodelcars.com      1      1143
Sales Rep
4      x3291  ljennings@classicmodelcars.com      1      1143
Sales Rep
..      ...      ...      ...      ...
...
95     x2312  mgerard@classicmodelcars.com        4      1102
Sales Rep
96     x2312  mgerard@classicmodelcars.com        4      1102
Sales Rep
97     x2312  mgerard@classicmodelcars.com        4      1102
Sales Rep
98     x2312  mgerard@classicmodelcars.com        4      1102
Sales Rep
99     x2312  mgerard@classicmodelcars.com        4      1102
Sales Rep

[100 rows x 21 columns]

```

---

## Expected Output

---

Notice that this also returned both columns: `salesRepEmployeeNumber` and `employeeNumber`. These columns contain identical values so you would probably actually only want to select one or the other.

<-brian-added/>

Finally, it is important to close the connection.

```
conn.close()
```

## Summary

In this lesson, you investigated joins. This included implementing the `ON` and `USING` clauses, aliasing table names, implementing `LEFT JOIN`, and using primary vs. foreign keys.